


## NoSQL & ACID DBMS – Common Misconceptions

- ACID Compliant DBMS don't support JSON and KVP
- Traditional DBMS can't handle the volume or the speed
- Web 2.0 Applications rely on JSON – traditional DBMS focus on text, integer, etc.
- Traditional DBMS don't work well with Web 2.0 development languages

8/18/2017 3:14 PM - Screen Clipping

# Postgres – NoSQL for the Enterprise

- Data Types
  - Postgres has JSON, JSONB, Key-Value Pair, plus arrays, ranges, timezones, dates, integer, floating point, etc.
- Performance Benchmarks
  - Postgres is very fast and can handle huge amounts of data
  - Postgres can selectively relax key ACID features to increase performance
- Proven track record
  - ACID compliant
  - Open source
  - ANSI SQL
  - Large developer and vendor community



8/18/2017 3:17 PM - Screen Clipping

# NoSQL Data in Postgres

- HSTORE

- Key-value pair
- Simple, fast and easy
- Postgres v 8.2 – pre-dates many NoSQL-only solutions

- JSON

- Hierarchical document model
- Introduced in Postgres 9.2, perfected in 9.3

- JSONB


- Binary version of JSON
- Faster, more operators and even more robust
- Postgres 9.4



8/18/2017 3:18 PM - Screen Clipping

# Postgres: Document Store

- JSON is the most popular data-interchange format on the web
- Derived from the ECMAScript Programming Language Standard (European Computer Manufacturers Association).
- Supported by virtually every programming language
- New supporting technologies continue to expand JSON's utility
  - PL/V8 JavaScript extension
  - Node.js
- Postgres native JSON data type (v9.2) and a JSON parser and a variety of JSON functions (v9.3)
- Postgres JSONB data type with binary storage and indexing (v9.4)



8/18/2017 3:19 PM - Screen Clipping

# Why JSON

- Wherever there is JavaScript you will find JSON
- Most languages support it
- Node.js is becoming popular.
- Lighter and more compact than XML.
- Most application don't need the rich structure of XML
- Flexible Structure
- Due to its flexible structure, JSON is a good fit for NoSQL.

8/18/2017 3:20 PM - Screen Clipping

# JSON Examples

- Creating a table with a JSONB field

```
CREATE TABLE json_data (data JSONB);
```

- Simple JSON data element:

```
{"name": "Apple Phone", "type": "phone", "brand":  
"ACME", "price": 200, "available": true,  
"warranty_years": 1}
```

- Inserting this data element into the table json\_data

```
INSERT INTO json_data (data) VALUES  
(' { "name": "Apple Phone",  
      "type": "phone",  
      "brand": "ACME",  
      "price": 200,  
      "available": true,  
      "warranty_years": 1  
} ')
```

# A simple query for JSON data

```
SELECT DISTINCT
    data->>'name' as products
FROM json_data;
```

products

-----

Cable TV Basic Service Package  
AC3 Case Black  
Phone Service Basic Plan  
AC3 Phone  
AC3 Case Green  
Phone Service Family Plan  
AC3 Case Red  
AC7 Phone

This query does not return JSON data – it returns text values associated with the key 'name'

# A query that returns JSON data

```
SELECT data FROM json_data;
```

```
data
```

```
-----  
  
{"name": "Apple Phone", "type": "phone",  
"brand": "ACME", "price": 200,  
"available": true, "warranty_years": 1}
```

This query returns the JSON data in its original format



# JSON Data Types

JSON is defined per RFC – 7159  
For more detail please refer  
<http://tools.ietf.org/html/rfc7159>

- 1. Number:
  - Signed decimal number that may contain a fractional part and may use exponential notation.
  - No distinction between integer and floating-point
- 2. String
  - A sequence of zero or more Unicode characters.
  - Strings are delimited with double-quotation mark
  - Supports a backslash escaping syntax.
- 3. Boolean
  - Either of the values true or false.
- 4. Array
  - An ordered list of zero or more values,
  - Each values may be of any type.
  - Arrays use square bracket notation with elements being comma-separated.
- 5. Object
  - An unordered associative array (name/value pairs).
  - Objects are delimited with curly brackets
  - Commas to separate each pair
  - Each pair the colon ':' character separates the key or name from its value.
  - All keys must be strings and should be distinct from each other within that object.
- 6. null
  - An empty value, using the word null

# JSON Data Type Example

```
{
  "firstName": "John",           -- String Type
  "lastName": "Smith",          -- String Type
  "isAlive": true,              -- Boolean Type
  "age": 25,                     -- Number Type
  "height_cm": 167.6,           -- Number Type
  "address": {                  -- Object Type
    "streetAddress": "21 2nd Street",
    "city": "New York",
    "state": "NY",
    "postalCode": "10021-3100"
  }
  "phoneNumbers": [             -- Object Array
    {                             -- Object
      "type": "home",
      "number": "212 555-1234"
    },
    {
      "type": "office",
      "number": "646 555-4567"
    }
  ],
  "children": [],
  "spouse": null                -- Null
}
```

8/18/2017 3:25 PM - Screen Clipping

# JSON and BSON

- BSON – stands for ‘Binary JSON’
- BSON != JSONB
  - BSON cannot represent an integer or floating-point number with more than 64 bits of precision.
  - JSONB can represent arbitrary JSON values.
- Caveat Emptor!
  - This limitation will not be obvious during early stages of a project!



8/18/2017 3:27 PM - Screen Clipping

# JSONB and Node.js - Easy as TT

```
// require the Postgres connector
var pg = require("pg");

// connection to local database
var conString = "pg://postgres:password@localhost:5432/nodetraining";

var client = new pg.Client(conString);
client.connect();

// initiate the sample database
client.query("CREATE TABLE IF NOT EXISTS emps(data jsonb)");
client.query("TRUNCATE TABLE emps;");
client.query('INSERT INTO emps VALUES($JSON$ {"firstname": "Ronald" , "lastname": "McDonald" }$JSON$)');
client.query('INSERT INTO emps values($JSON$ {"firstname": "Mayor", "lastname": "McCheese"}$JSON$)');

// run SELECT query
client.query("SELECT * FROM emps",function(err,result){
    console.log("Test Output of JSON Result Object");
    console.log(result);
    console.log("Parsed rows");

    // parse the result set
    for (var i = 0; i< result.rows.length ; i++){
        var data = JSON.parse(result.rows[i].data);
        console.log("First Name => " + data.firstname + "\t| Last Name => " + data.lastname);
    }
    client.end();
})|
```

8/18/2017 3:29 PM - Screen Clipping

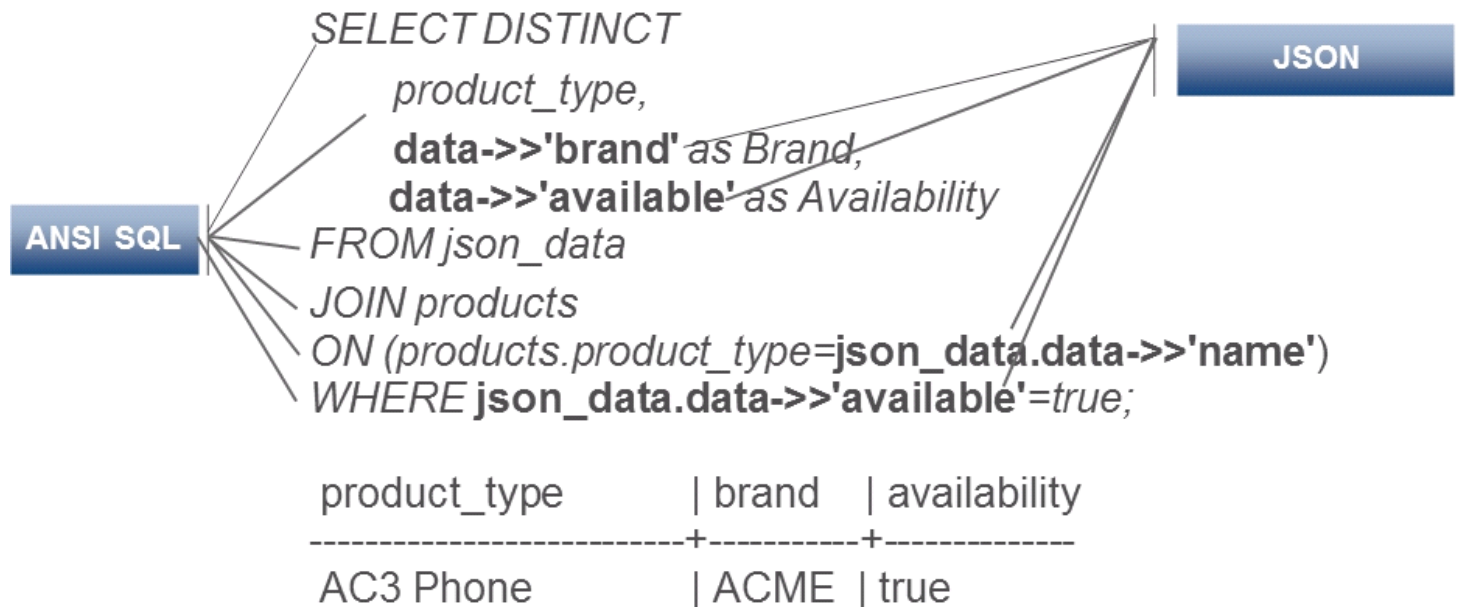
# JSON and ANSI SQL – A Great Fit



- JSON is naturally integrated with ANSI SQL in Postgres
- JSON and HSTORE are elegant and easy to use extensions of the underlying object-relational model
- JSON and SQL queries use the same language, the same planner, and the same ACID compliant transaction framework

8/18/2017 3:30 PM - Screen Clipping

# JSON and ANSI SQL Example



No need for programmatic logic to combine SQL and NoSQL in the application – Postgres does it all



# Bridging between SQL and NoSQL

## Simple ANSI SQL Table Definition

```
CREATE TABLE products (id integer, product_name text );
```

## Select query returning standard data set

```
SELECT * FROM products;
```


```
id | product_name
----+-----
1  | iPhone
2  | Samsung
3  | Nokia
```

## Select query returning the same result as a JSON data set

```
SELECT ROW_TO_JSON(products) FROM products;
```

```
{"id":1,"product_name":"iPhone"}
{"id":2,"product_name":"Samsung"}
{"id":3,"product_name":"Nokia"}
```

# Postgres Provides Great Flexibility



- Start unstructured, and become structured as you learn more
  - Use the quick-to-get-started capabilities of NoSQL
  - Complete the initial sprints without a DBA
  - Move data between unstructured and structured
  - Embrace corporate data standards as you move from the stand-alone application towards integrated applications with a bigger value proposition

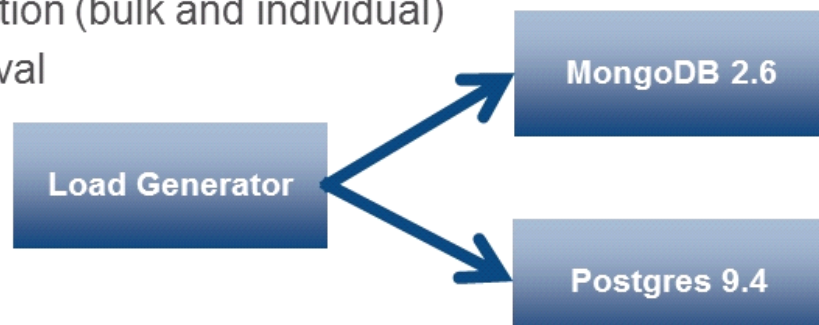
**By 2017, 50% of data stored in NoSQL DBMSs will be damaging to the business due to lack of applied information governance policies and programs.**

Gartner, December 2013




# Postgres NoSQL Performance Evaluation

- Goal
  - Help our customers understand when to choose Postgres and when to choose a specialty solution
  - Help us understand where the NoSQL limits of Postgres are
- Setup
  - Compare Postgres 9.4 to Mongo 2.6
  - Single instance setup on AWS M3.2XLARGE (32GB)
- Test Focus
  - Data ingestion (bulk and individual)
  - Data retrieval



8/18/2017 3:36 PM - Screen Clipping

# NoSQL Performance Evaluation



|               | Postgres | MongoDB |
|---------------|----------|---------|
| Data Load (s) | 4,732    | 13,046  |
| Insert (s)    | 29,236   | 86,253  |
| Select (s)    | 594      | 2,763   |
| Size (GB)     | 69       | 145     |

## Correction to earlier versions:

MongoDB console does not allow for INSERT of documents > 4K. This lead to truncation of the MongoDB size by approx. 25% of all records in the benchmark.

# Postgres NoSQL Performance Evaluation

- Tests confirm that Postgres can handle many NoSQL workloads
- EDB is making the test scripts publically available
- EDB encourages community participation to better define where Postgres should be used and where specialty solutions are appropriate
- Download the source at [https://github.com/EnterpriseDB/pg\\_nosql\\_benchmark](https://github.com/EnterpriseDB/pg_nosql_benchmark)



8/18/2017 5:16 PM - Screen Clipping

## Foreign Data Wrappers – Co-Existence Platform

- FDW implements SQL/MED ("SQL Management of External Data")
- PostgreSQL 9.1 - read-only support
- PostgreSQL 9.3 – read/write support
- FDW
  - Makes data on other servers (or services) look like tables in Postgres
  - available for databases (MongoDB, MySQL, Oracle, ...), files, services (Twitter, ...)
- MongoDB FDW: <https://github.com/EnterpriseDB>



# MongoDB FDW Example

```
CREATE EXTENSION mongo_fdw;

CREATE SERVER mongo_server
    FOREIGN DATA WRAPPER mongo_fdw
    OPTIONS (address '172.24.39.129', port '27017');

CREATE USER MAPPING FOR enterprisedb
    SERVER mongo_server
    OPTIONS (username 'mongo', password 'mongo');

CREATE FOREIGN TABLE mongo_data(
    name text,
    brand text,
    type text)
    SERVER mongo_server
    OPTIONS (
        database 'benchmark',
        collection 'json_tables');
```

# MongoDB FDW Example

```
SELECT * FROM mongo_data WHERE brand='ACME' limit 10;
```

| name         | brand | type  |
|--------------|-------|-------|
| AC7553 Phone | ACME  | phone |
| AC7551 Phone | ACME  | phone |
| AC7519 Phone | ACME  | phone |
| AC7565 Phone | ACME  | phone |
| AC7555 Phone | ACME  | phone |
| AC7529 Phone | ACME  | phone |
| AC7528 Phone | ACME  | phone |
| AC7547 Phone | ACME  | phone |
| AC7587 Phone | ACME  | phone |
| AC7541 Phone | ACME  | phone |

(10 rows)

8/18/2017 5:20 PM - Screen Clipping

# MongoDB FDW Example

```
INSERT INTO mongo_data(name, brand, type)
VALUES('iphone6 phone','Apple Inc','phone');
```

```
SELECT* FROM mongo_data WHERE brand='Apple Inc';
```

| _id                      | name          | brand     | type  |
|--------------------------|---------------|-----------|-------|
| 53ea4f59fe5586a15714881d | iphone6 phone | Apple Inc | phone |

```
UPDATE mongo_data SET brand='Apple Product'
WHERE brand='Apple Inc';
```

```
SELECT * FROM mongo_data WHERE brand='Apple Product';
```

| _id                      | name          | brand         | type  |
|--------------------------|---------------|---------------|-------|
| 53ea4f59fe5586a15714881d | iphone6 phone | Apple Product | phone |

8/18/2017 5:24 PM - Screen Clipping

# “No SQL Only” or “Not Only SQL”?

- Structures and standards emerge!
- Data has references (products link to catalogues; products have bills of material; components appear in multiple products; storage locations link to ISO country tables)
- When the database has duplicate data entries, then the application has to manage updates in multiple places – what happens when there is no ACID transactional model?

8/18/2017 5:26 PM - Screen Clipping

- NoSQL only databases are not ACID transactional. Which means, if you tell go and change this data element in 5 collections, there is no guarantee to check if all of your changes happen or none of your changes happen.
- There is no concept of atomicity in NoSQL only databases

# Postgres: The Best of Both Worlds

- Postgres has many NoSQL features without the drawbacks:
  - Schema-less data combined with structured data
  - High performance with predictable transaction model
  - Durable by default, but configurable per-table or per-transaction
  - Standards based with very low technology risk
  - Foreign Data Wrappers (FDW) for co-existence
  - Highly available skill set

8/18/2017 5:33 PM - Screen Clipping



# Postgres: The Best of Both Worlds

- Postgres has many NoSQL features without the drawbacks:
  - Schema-less data combined with structured data
  - High performance with predictable transaction model
  - Durable by default, but configurable per-table or per-transaction
  - Standards based with very low technology risk
  - Foreign Data Wrappers (FDW) for co-existence
  - Highly available skill set

8/18/2017 5:35 PM - Screen Clipping

# Say 'Yes' to 'Not Only SQL'

- Postgres is Not Only SQL (NoSQL is No SQL only)
- Fully ACID compliant
- Proven track record
- Fully capable of handling the variety, velocity and volume requirements of most applications
- Tackle NoSQL projects without leaving the capabilities of the relational model behind you
- Combine Oracle compatibility, JSON and PostGIS to migrate applications onto more cost-effective platforms, make the app NoSQL capable and geo-location aware.

8/18/2017 5:36 PM - Screen Clipping